

## CO4 -HOME ASSIGNMENT

1. Compare and contrast the characteristics of key-value, document, and graph NoSQL databases.

**Key-Value Databases**

- **Model:** Simple key-value pairs.
- **Schema:** Schema-less.
- **Performance:** Fast lookups by key.
- **Querying:** Limited to key-based lookups.
- **Use Case:** Caching, session storage.
- **Examples:** Redis, DynamoDB.

**Document Databases**

- **Model:** JSON/BSON documents.
- **Schema:** Flexible, semi-structured.
- **Performance:** Good for nested data.
- **Querying:** Rich querying, filtering, aggregation.
- **Use Case:** Content management, user profiles.
- **Examples:** MongoDB, CouchDB.

**Graph Databases**

- **Model:** Nodes and edges.
- **Schema:** Flexible, relationships first-class.
- **Performance:** Optimized for connected data.
- **Querying:** Graph traversal languages.
- **Use Case:** Social networks, fraud detection.
- **Examples:** Neo4j, ArangoDB.

2. Describe structured, unstructured, and semi-structured data types in bigdata.

## Structured Data

- Organized in rows and columns.
- Stored in RDBMS.
- Easy to query (SQL).
- *Example:* Excel sheets, SQL tables.

## Unstructured Data

- No fixed format.
- Hard to process directly.
- *Example:* Images, videos, emails.

## Semi-Structured Data

- Partial structure (tags, keys).
- Not in tables, but organized.
- *Example:* JSON, XML, log files.

3. Consider the following simplified schema for an airline reservation system: { "\_id": ObjectId, "flightNumber": String, "departure": { "airport": String, "city": String, "time": Date }, "arrival": { "airport": String, "city": String, "time": Date }, "seatsAvailable": Number } i. Create a new reservation for a passenger on a specific flight. ii. Find all flights departing from a specific airport. iii. Find all passengers with a specific passport number. iv. Find all reservations for a specific flight. v. Update the contact details of a specific passenger.

**i. Create a new reservation for a passenger on a specific flight**

```
db.reservations.insertOne({  
  flightNumber: "AI101",  
  passengerId: ObjectId("PASSENGER_OBJECT_ID"),  
  reservationDate: new Date(),  
  seatNumber: "12A"  
});
```

---

**ii. Find all flights departing from a specific airport**

```
db.flights.find({ "departure.airport": "JFK" });
```

---

**iii. Find all passengers with a specific passport number**

```
db.passengers.find({ passportNumber: "A1234567" });
```

---

**iv. Find all reservations for a specific flight**

```
db.reservations.find({ flightNumber: "AI101" });
```

---

**v. Update the contact details of a specific passenger**

```
db.passengers.updateOne(  
  { _id: ObjectId("PASSENGER_OBJECT_ID") },  
  { $set: { contact: { phone: "9876543210", email: "newemail@example.com" } } }  
);
```

#### 4. Explain the CAP theorem in NOSQL databases.

### **CAP Theorem (NoSQL)**

In a distributed system, you can only guarantee 2 out of 3:

- **C - Consistency:** All nodes show the same data.
- **A - Availability:** Every request gets a response.
- **P - Partition Tolerance:** Works despite network failures.

No system can guarantee all 3.

- **CP:** Consistency + Partition (e.g., MongoDB)
- **AP:** Availability + Partition (e.g., DynamoDB)
- **CA:** Not possible if there's a network failure.

5. Describe the role of YARN (Yet Another Resource Negotiator) in Hadoop architecture. How does it manage resources and schedule jobs in a cluster? What are the architectural differences and advantages of each?

## YARN in Hadoop

- Manages **resources** and **job scheduling** in Hadoop.
- Allows multiple data engines (MapReduce, Spark) to run on the same cluster.

## Key Components

- **ResourceManager (RM)**: Allocates resources.
- **NodeManager (NM)**: Manages each node's resources.
- **ApplicationMaster (AM)**: Manages each app's execution.
- **Containers**: Run tasks using allocated resources.

## Differences from Classic MapReduce

- Replaces JobTracker with RM + AM.
- Supports multiple processing models.
- More scalable and flexible.

## Advantages

- Better resource use.
- Runs multiple apps.
- Scalable and efficient.

## 6. Explain mapreduce programming in Hadoop.

### MapReduce in Hadoop

MapReduce is a programming model for processing large datasets in parallel.

#### 1. Map Phase:

- **Input:** Key-value pairs.
- **Operation:** Mapper processes each record and produces intermediate key-value pairs.

#### 2. Shuffle & Sort:

- Groups the same keys together.

#### 3. Reduce Phase:

- **Input:** Grouped key-value pairs.
- **Operation:** Reducer aggregates the values (e.g., summing).

### Advantages

- Scalable and parallel processing.
- Fault-tolerant.