

# DEPARTMENT OF CSE-H

## TOPIC INTRODUCTION MONGODB

---

# WHAT IS MONGODB?

- Developed by 10gen
  - Founded in 2007
- A document-oriented, NoSQL database
  - Hash-based, *schema-less database*
    - No Data Definition Language
    - In practice, this means you can store hashes with any keys and values that you choose
      - Keys are a basic data type but in reality stored as strings
      - Document Identifiers (\_id) will be created for each document, field name reserved by system
    - Application tracks the schema and mapping
    - Uses BSON format
      - Based on JSON – B stands for Binary
- Written in C++
- Supports APIs (drivers) in many computer languages
  - JavaScript, Python, Ruby, Perl, Java, Java Scala, C#, C++, Haskell, Erlang

# FUNCTIONALITY OF MONGODB

- Dynamic schema
  - No DDL
- Document-based database
- Secondary indexes
- Query language via an API
- Atomic writes and fully-consistent reads
  - If system configured that way
- Master-slave replication with automated failover (replica sets)
- Built-in horizontal scaling via automated range-based partitioning of data (sharding)
- No joins nor transactions

# WHY USE MONGODB?

---

- Simple queries
- Functionality provided applicable to most web applications
- Easy and fast integration of data
  - No ERD diagram
- Not well suited for heavy and complex transactions systems

## SESSION INTRODUCTION

# RDB CONCEPTS TO NO SQL

RDBMS		MongoDB
Database		Database
Table, View	→ →	Collection
Row	→	Document (BSON)
Column		Field
Index	→	Index
Join	→	Embedded Document
Foreign Key	→ →	Reference
Partition	→	Shard

Collection is not strict about what it Stores

Schema-less

Hierarchy is evident in the design

Embedded Document

## SESSION INTRODUCTION

# MONGODB PROCESSES AND CONFIGURATION

- Mongod – Database instance
- Mongos - Sharding processes
  - Analogous to a database router.
  - Processes all requests
  - Decides how many and which *mongods* should receive the query
  - *Mongos* collates the results, and sends it back to the client.
- Mongo – an interactive shell ( a client)
  - Fully functional JavaScript environment for use with a MongoDB
- You can have one *mongos* for the whole system no matter how many *mongods* you have
- OR you can have one local *mongos* for every client if you wanted to minimize network latency.

# DATABASES AND COLLECTIONS

MongoDB stores data records as documents (specifically BSON documents) which are gathered together in collections. A database stores one or more collections of documents.

## Create a Database

If a database does not exist, MongoDB creates the database when you first store data for that database.

### Syntax

Basic syntax of use DATABASE statement is as follows –  
`use DATABASE_NAME`

### The dropDatabase() Method

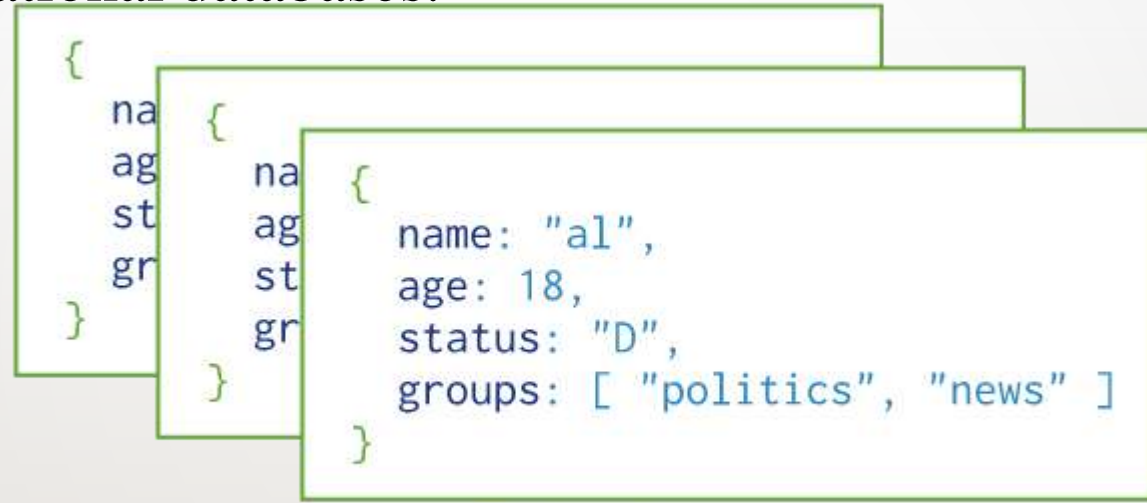
MongoDB `db.dropDatabase()` command is used to drop a existing database.

### Syntax

Basic syntax of `dropDatabase()` command is as follows –  
`db.dropDatabase()`

## Collections

MongoDB stores documents in collections. Collections are analogous to tables in relational databases.



Collection

The createCollection() Method

MongoDB **db.createCollection(name)** is used to create collection.

Syntax

Basic syntax of **createCollection()** command is as follows –

**db.createCollection(name)**



## MongoDB supports many datatypes.

- String: A basic and commonly used data type
- Integer: Used to store numeric values that don't require decimal precision
- Double: Stores numeric numbers with 8 bytes (64-bit IEEE 754 floating point)
- Boolean: Stores true or false values
- Array: Stores arrays
- Object: Stores embedded documents as key-value pairs within other documents
- Date: Stores the current date or time
- ObjectId: A 12-byte identifier that uniquely identifies documents within a collection

# BSON FORMAT


---

- Binary-encoded serialization of JSON-like documents
- Zero or more key/value pairs are stored as a single entity
- Each entry consists of a field name, a data type, and a value
- Large elements in a BSON document are prefixed with a
- length field to facilitate scanning

## SESSION INTRODUCTION

# SCHEMA FREE

- MongoDB does not need any pre-defined data schema
- Every document in a collection could have different data
  - Addresses NULL data fields



```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la ciacco"],  
  loc: [32.7, 63.4],  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  loc: [40.7, 73.4],  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  loc: [44.6, 71.3]}
```

# JSON FORMAT

- Data is in name / value pairs
- A name/value pair consists of a field name followed by a colon, followed by a value:
  - Example: “name”: “R2-D2”
- Data is separated by commas
  - Example: “name”: “R2-D2”, race : “Droid”
- Curly braces hold objects
  - Example: {“name”: “R2-D2”, race : “Droid”, affiliation: “rebels”}
- An array is stored in brackets []
  - Example [ {“name”: “R2-D2”, race : “Droid”, affiliation: “rebels”}, {“name”: “Yoda”, affiliation: “rebels”} ]

# MONGODB FEATURES

- Document-Oriented storage
- Full Index Support
- Replication & High Availability
- Auto-Sharding
- Querying
- Fast In-Place Updates
- Map/Reduce functionality

Agile

Scalable

# CRUD OPERATIONS

- Create
  - `db.collection.insert( <document> )`
  - `db.collection.insertOne( <document> )`
  - `db.collection.insertMany( [<document>, [<document>, ..] )`
- Read
  - `db.collection.find( <query>, <projection> )`
  - `db.collection.findOne( <query>, <projection> )`
- Update
  - `db.collection.update( <query>, <update>, <options> )`
  - `db.collection.updateOne( <query>, <update>, { upsert: true } )`
  - `db.collection.updateMany( <query>, <update>, { upsert: true } )`
- Delete
  - `db.collection.deleteMany( <query> )`
  - `db.collection.deleteOne( <query> )`

Collection specifies the collection or the  
'table' to store the document

# MONGODB - INSERT DOCUMENT

## Insert Documents

There are 2 methods to insert documents into a MongoDB database.

### `insertOne()`

To insert a single document, use the `insertOne()` method.

This method inserts a single object into the database.

### Syntax

- The basic syntax of insert() command is as follows –
- `>db.COLLECTION_NAME.insert(document)`

Db.collection specifies the collection or the ‘table’ to store the document

- `db.collection_name.insert( <document> )`
  - Omit the `_id` field to have MongoDB generate a unique key
  - Example `db.parts.insert( {type: “screwdriver”, quantity: 15 } )`
  - `db.parts.insert({_id: 10, type: “hammer”, quantity: 1 } )`

## insertMany()

To insert multiple documents at once, use the **insertMany()** method.  
This method inserts an array of objects into the database.

### Syntax:

```
db.Collection_name.insertMany(  
[<document 1>, <document 2>, ...])
```



# MONGODB – FIND OPERATIONS

**MongoDB** provides powerful methods for retrieving documents from its collections with **find()** and **findOne()** methods. The **find()** method supports **complex queries** with various **operators** and allows specifying which fields to include or exclude and optimizing performance by using indexes.

## Find() Method

- It is a **primary** method for retrieving documents from a collection in [MongoDB](#).
- It Supports a wide range of query operators, including [comparison](#), [logical](#) and **element** operators also allowing for complex queries.
- It allows us to specify which fields to include or exclude in the result set and reduce the amount of data transferred.
- It automatically uses **indexes** to optimize query performance also improving speed and efficiency.

## Syntax:

```
db.Collection_name.find(selection_criteria, projection,options)
```

## findOne() Method

- findOne() in MongoDB is a method used to retrieve a single document from a collection that matches a specified criteria.
- It returns **only one document**, even if multiple documents match the criteria.
- If no matching document is found, it returns null.
- You can optionally provide a **query criteria** to specify which document to find.

# MONGODB – UPDATE() METHOD

**MongoDB update operations** allow us to modify documents in a collection. These operations can update a single document or multiple documents based on specified criteria. MongoDB offers various update operators to perform specific actions like **setting a value, incrementing a value** or **updating elements within arrays**.

MongoDB update()

The MongoDB update() method is a method that is used to update a single document or multiple documents in the collection. When the document is updated the \_id field remains unchanged.

The db.collection.update() method updates a single document by default. To update all documents that match the given query, use the multi: true option. Include the option “multi: true”.

## Syntax

```
db.COLLECTION_NAME.update({SELECTI  
ON_CRITERIA},  
{ $set: {UPDATED_DATA} }, {  
  upsert: <boolean>,  
  multi: <boolean>,  
  writeConcern: <document>,  
  collation: <document>,  
  arrayFilters: [ <filterdocument1>,  
    ... ],  
  hint: <document|string>  
})
```

# MONGODB – UPDATE() METHOD

## MongoDB Update() Method

---

The update() method updates the values in the existing document.

### Syntax

The basic syntax of **update()** method is as follows –

```
>db.COLLECTION_NAME.update(SELECTION_CRITERIA,  
UPDATED_DATA)
```

- `db.collection_name.update( <query>, <update>, { upsert: true } )`
  - Will update 1 or more records in a collection satisfying query
- `db.collection_name.findAndModify(<query>, <sort>, <update>,<new>, <fields>,<upsert>)`
  - Modify existing record(s) – retrieve old or new version of the record

# DELETE OPERATIONS

## `db.collection.deleteOne()`

`db.collection.deleteOne()` deletes the first document that matches the filter. Use a field that is part of a unique index such as `_id` for precise deletions.

`db.collection.deleteMany()` operation successfully deletes one or more documents

- `db.collection_name.remove(<query>, <justone>)`
  - Delete all records from a collection or matching a criterion
  - `<justone>` - specifies to delete only 1 record matching the criterion
  - Example: `db.parts.remove(type: /^h/ }` ) - remove all parts starting with h
  - `Db.parts.remove()` – delete all documents in the parts collections

# MONGODB PROJECTION

- In MongoDB, projection means selecting only the necessary data rather than selecting whole of the data of a document. If a document has 5 fields and you need to show only 3, then select only 3 fields from them.

## Syntax

The basic syntax of **find()** method with projection is as follows

```
>db.COLLECTION_NAME.find({}, {KEY:1})
```

## Limit() Method

To limit the records in MongoDB, you need to use **limit()** method. The method accepts one number type argument, which is the number of documents that you want to be displayed.

## Syntax

The basic syntax of **limit()** method is as follows –

```
>db.COLLECTION_NAME.find().limit(NUMBER)
```

## MongoDB Skip() Method

Apart from limit() method, there is one more method **skip()** which also accepts number type argument and is used to skip the number of documents.

### Syntax

The basic syntax of **skip()** method is as follows –

```
>db.COLLECTION_NAME.find().limit(NUMBER).skip(NUMBER)
```

## Example

Following example will display only the second document.

```
>db.mycol.find({}, {"title":1, _id:0}).limit(1).skip(1)
{"title":"NoSQL Overview"} >
```

## MongoDB – sort() Method

- To sort documents in MongoDB, you need to use **sort()** method. The method accepts a document containing a list of fields along with their sorting order. To specify sorting order 1 and -1 are used. 1 is used for ascending order while -1 is used for descending order.

### Syntax

The basic syntax of **sort()** method is as follows –

```
>db.COLLECTION_NAME.find().sort({KEY:1})
```



- MongoDB count()
- The count() method in MongoDB is a simple and effective way to count the number of documents that meet specific criteria in a collection. It can take an optional query parameter to filter the documents before counting. The count() method is available in both the db.collection and cursor objects in MongoDB.
- db.collection.count() or db.collection.countDocuments()
- Returns the count of documents that would match a find() query for the collection or view. The db.collection.count() method does not perform the find() operation but instead counts and returns the number of results that match a query.

# MONGODB – FINDANDMODIFY() METHOD

---

The **MongoDB** findAndModify method is used to modify and return a **single document** that matches the given criteria. It can be used to **update**, **remove** or **insert** documents atomically.

```
db.collection.findOneAndDelete( filter, options )
```

# MONGODB – FINDONE AND DELETE() METHOD

---

The **MongoDB findOneAndDelete()** is a powerful tool for deleting documents from a collection based on specific criteria. It allows for the deletion of a single document that matches the query filter.

## Syntax:

```
db.Collection_name.findOneAndDelete(  
  Selection_criteria,  
  {  
    projection: <document>,  
    sort: <document>,  
    maxTimeMS: <number>,  
    collation: <document>  
  })
```

# MONGODB FINDONE AND REPLACE() METHOD

- 
- The **findOneAndReplace()** method replaces the first matched document based on the selection criteria.  
replacing
  - By default, this method returns the original document. To return the replacement document, set the value of the returnNewDocument option to true.
  - **db.collection.findOneAndReplace( filter, replacement, options )**

## QUERY OPERATORS

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to a specified value
\$lt, \$lte	Matches values less than or ( equal to ) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

## CRUD EXAMPLES

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.find (  
  { "_id" : ObjectId("51"),  
    "first" : "John",  
    "last" : "Doe",  
    "age" : 39  
  }  
)
```

```
> db.user.update(  
  { "_id" : ObjectId("51") },  
  {  
    $set: {  
      age: 40,  
      salary: 7000  
    }  
  }  
)
```

```
> db.user.remove(  
  {  
    "first": /^J/  
  }  
)
```

# SQL VS. MONGO DB ENTITIES

My SQL	Mongo DB
<pre> START TRANSACTION; INSERT INTO contacts VALUES   (NULL, 'joeblow'); INSERT INTO contact_emails VALUES   ( NULL, "joe@blow.com",     LAST_INSERT_ID() ),   ( NULL,     "joseph@blow.com",     LAST_INSERT_ID() ); COMMIT; </pre>	<pre> db.contacts.save( {   userName: "joeblow",   emailAddresses: [     "joe@blow.com",     "joseph@blow.com" ] } ); </pre> <p>Similar to IDS from the 70's        Bachman's brainchild</p> <p>DIFFERENCE:        MongoDB separates physical structure        from logical structure</p> <p>Designed to deal with large &amp; distributed</p>

37

---

## Reference Books:

1. [NoSQL with MongoDB in 24 Hours, Sams Teach Yourself 1st Edition](#) by Brad Dayley
2. [MongoDB Basics 1st Edition](#) by Peter Membrey, David Hows, and Eelco Plugge

## Sites and Web links:

1. <https://www.mongodb.com/docs/manual/faq/>